

Understanding Dimension, Shape, Axis, Broadcasting and Vectorization

1. Dimension (ndim)

- **Definition:** Number of axes (or levels of nesting) in a NumPy array.
- **Examples:**
 - Scalar → 0-D (No dimension)
 - Vector → 1-D
 - Matrix → 2-D
 - Tensor → 3-D or higher

| Example Array | ndim (Dimension) |
|----------------------------------|------------------|
| <code>np.array(5)</code> | 0-D |
| <code>np.array([1,2,3])</code> | 1-D |
| <code>np.array([[1,2,3]])</code> | 2-D |
| <code>np.zeros((2,3,4))</code> | 3-D |

2. Shape

- **Definition:** A tuple representing the **size of the array along each dimension**.
- **Examples:**
 - A 1-D array of 5 elements → `(5,)`
 - A 2-D array with 3 rows and 4 columns → `(3, 4)`
 - A 3-D array with shape `(2, 3, 4)` → 2 blocks, 3 rows, 4 columns each

3. Axis

- **Definition:** A particular **dimension along which operations are performed**.

Rule of Axes in NumPy

Axis number increases with depth of dimensions:

- Axis 0 → First dimension (outermost)
- Axis 1 → Second dimension
- Axis 2 → Third dimension ...and so on.

Operations collapse the given axis, meaning elements along that axis are combined, and the axis disappears in the result.

- **Example:**

- Summing along `axis=0` → Collapse first dimension
- Summing along `axis=1` → Collapse second dimension

1. 2D Array Case

```
In [4]: import numpy as np
arr = np.array([[1, 2, 3],
                [4, 5, 6]])
```

```
In [5]: np.sum(arr, axis=0) # Sum along axis 0 (Column-wise)
```

```
Out[5]: array([5, 7, 9])
```

```
In [6]: np.sum(arr, axis=1) # Sum along axis 1 (Row-wise)
```

```
Out[6]: array([ 6, 15])
```

2. 3D Array Case

```
In [8]: import numpy as np

arr3d = np.array([
    [[ 1,  2,  3],
     [ 4,  5,  6]],

    [[ 7,  8,  9],
     [10, 11, 12]]
])

arr3d
```

```
Out[8]: array([[[ 1,  2,  3],
                 [ 4,  5,  6]],

                 [[ 7,  8,  9],
                  [10, 11, 12]]])
```

```
In [9]: print(arr3d.shape)
```

```
(2, 2, 3)
```

Shape (2, 2, 3) means:

- 2 blocks
- 2 rows per block
- 3 columns per row

```
In [11]: np.sum(arr3d, axis=0)
```

```
Out[11]: array([[ 8, 10, 12],
 [14, 16, 18]])
```

```
In [14]: np.sum(arr3d, axis=0).shape
```

```
Out[14]: (2, 3)
```

```
In [12]: np.sum(arr3d, axis=1)
```

```
Out[12]: array([[ 5,  7,  9],
 [17, 19, 21]])
```

```
In [13]: np.sum(arr3d, axis=1).shape
```

```
Out[13]: (2, 3)
```

```
In [15]: np.sum(arr3d, axis=2)
```

```
Out[15]: array([[ 6, 15],
 [24, 33]])
```

```
In [16]: np.sum(arr3d, axis=2).shape
```

```
Out[16]: (2, 2)
```

In Short:

- **Dimension (ndim)** → How many axes?
 - **Shape** → Size along each axis.
 - **Axis** → Direction along which operations are applied.
-

4. Broadcasting

Broadcasting Rules

1. Compare shapes from the last dimension to the first.
 - If one shape has fewer dimensions, leading 1s are added automatically to the smaller shape.
2. Dimensions are compatible if:
 - They are equal, or
 - One of them is 1
3. Resulting shape is the maximum in each dimension after broadcasting.

```
In [19]: #Simple Example (1D + Scalar)
```

```
arr = np.array([1, 2, 3])
print(arr + 5) # Broadcasts 5 to [5, 5, 5]
```

```
[6 7 8]
```

```
In [20]: # 1D and 2D broadcasting
```

```
A = np.array([[1, 2, 3],
             [4, 5, 6]]) # Shape (2,3)

b = np.array([10, 20, 30]) # Shape (3,)

C = A + b
print(C)
```

```
[[11 22 33]
 [14 25 36]]
```

```
In [21]: np.array([1,2,3]).shape
```

```
Out[21]: (3,)
```

```
In [22]: np.array([[1],[2]]).shape
```

```
Out[22]: (2, 1)
```

```
In [23]: np.array([1,2,3]) + np.array([[1],[2]])
```

```
Out[23]: array([[2, 3, 4],
                 [3, 4, 5]])
```

Broadcasting Fails When Shapes are Incompatible

```
In [26]: A = np.array([[1, 2, 3],
                     [4, 5, 6]]) # Shape (2,3)
```

```
B = np.array([[10, 20],
             [30, 40]]) # Shape (2,2)
```

```
# A+B
```

Broadcasting Internals Explained with Examples

1. Broadcasting in NumPy is implemented in C under the hood for speed.
2. Align shapes: `(2,3)` and `(1,3)`
3. Strides for B (row stride = 0) → Reuse the same row twice
4. Compute element-wise sum in a new `(2,3)` array - using Vectorization

Advantages of Broadcasting

1. Memory efficient (no actual data replication).

2. Faster computations using vectorization instead of loops.

5. Vectorization

- Vectorization means performing operations on entire arrays at once instead of iterating element by element in Python.
- NumPy operations are implemented in C, which avoids Python-level loops and interpreter overhead.

```
In [31]: A = np.array([1, 2, 3])
B = np.array([2, 3, 4])
print(A,B)
```

[1 2 3] [2 3 4]

```
In [32]: A+B
```

Out[32]: array([3, 5, 7])

```
In [41]: [int(x+y) for x, y in zip(A, B)]
```

Out[41]: [3, 5, 7]

When you do A + B:

- NumPy does not loop in Python.
- It calls optimized C functions that perform element-wise addition directly on the memory buffers.

Why Vectorization is Faster?

- **C-level implementation:** NumPy's loops are in compiled C, much faster than Python loops.
- **Fewer Python instructions:** No per-element Python function call overhead.
- **SIMD optimizations:** Uses CPU vectorized instructions for multiple elements at once using wide vector registers (SSE, AVX registers).
- **Memory efficiency:** Works with contiguous arrays and broadcasting without creating many intermediate Python objects.

Performance Comparison - Vectorization Vs Loops

```
In [ ]: # Example 1: Squaring 10 million numbers
```

```
In [43]: import numpy as np
import time
```

```
N = 10000000
arr = np.arange(N)

arr.shape
```

Out[43]: (10000000,)

```
In [44]: # Python Loop Implementation
start = time.time()
result_loop = [x**2 for x in arr] # Python Loop
print("Python loop time:", time.time() - start)
```

Python loop time: 1.0810484886169434

```
In [45]: # NumPy Vectorized Operation (Fast)
start = time.time()
result_vec = arr**2 # Vectorized operation
print("Vectorized time:", time.time() - start)
```

Vectorized time: 0.01813983917236328

```
In [46]: 1.0810484886169434// 0.01813983917236328
```

Out[46]: 59.0

⚡ Vectorized code is **50x+ faster** because it runs in compiled C loops internally.